

```

import pandas as pd

from google.colab import drive
drive.mount('/content/drive')

↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

import keras
print(keras.__version__)

↳ 3.5.0

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt

BATCH_SIZE=32
IMAGE_SIZE=640
CHANNELS=3
EPOCHS=50

data1=pd.read_csv('/content/drive/MyDrive/Disease dataset/test/_classes.csv')
data2=pd.read_csv('/content/drive/MyDrive/Disease dataset/train/_classes.csv')
data3=pd.read_csv('/content/drive/MyDrive/Disease dataset/valid/_classes.csv')

data1.head()
data2.head()
data3.head()

```

	filename	(BRD)	Bovine	Contagious	Dermatitis	Disease	Ecthym	Respiratory	Unlabeled	h
0	istockphoto-1319281522-612x612.jpg.rf.995226a1...	0	0	0	0	0	0	0	0	0
1	IMG_20220829_124626.jpg.rf.a354942b646265839d9...	0	0	0	0	0	0	0	0	0
2	images_jpeg.rf.a5131ece4904d8dcaf77e31b725837...	1	1	0	0	1	0	1	0	0
3	IMG-20220830-WA0065.jpg.rf.9b091d6fbdde75a3d8e...	0	0	0	0	0	0	0	0	0
4	Ayrshirecattle69.jpg.rf.983d9add00cdf1d9e36b3d...	0	0	0	0	0	0	0	0	0

Next steps: [Generate code with data3](#) [View recommended plots](#) [New interactive sheet](#)

```

dataset=tf.keras.preprocessing.image_dataset_from_directory(
    "/content/drive/MyDrive/Disease dataset",
    seed=123,
    shuffle=True,
    image_size=(IMAGE_SIZE,IMAGE_SIZE),
    batch_size=BATCH_SIZE
)

```

↳ Found 830 files belonging to 3 classes.

```

class_names=dataset.class_names
class_names

```

↳ ['test', 'train', 'valid']

```

for image_batch,labels_batch in dataset.take(1):
    print(image_batch.shape)
    print(labels_batch.numpy())
    for i in range(12):
        ax=plt.subplot(3,4,i+1)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[labels_batch[i]])
        plt.axis("off")

```

```
↳ (32, 640, 640, 3)
[1 2 2 0 1 1 2 1 1 1 1 1 1 2 2 0 1 1 2 1 1 1 1 2 1 1 0 1 2 2 1]
```



```
train_size=0.8
len(dataset)*train_size
```

```
↳ 20.8
```

```
train_ds=dataset.take(20)
len(train_ds)
```

```
↳ 20
```

```
test_ds=dataset.skip(20)
len(test_ds)
```

```
↳ 6
```

```
val_size=0.1
len(dataset)*val_size
```

```
↳ 2.6
```

```
val_ds=test_ds.take(2)
len(val_ds)
```

```
↳ 2
```

```
def get_dataset_partitions_tf(ds,train_split=0.8,val_split=0.1,test_split=0.1,shuffle=True,shuffle_size=10000):
    assert(train_split+test_split+val_split)==1
    ds_size=len(ds)
    train_size=int(train_split*ds_size)
    val_size=int(val_split*ds_size)
    test_size=int(test_split*ds_size)
    if shuffle:
        ds=ds.shuffle(shuffle_size,seed=12)
    train_ds=ds.take(train_size)
    val_ds=ds.skip(train_size).take(val_size)
    test_ds=ds.skip(train_size+val_size)
    return train_ds,val_ds,test_ds
```

```
train_ds,val_ds,test_ds=get_dataset_partitions_tf(dataset)
len(train_ds)
len(val_ds)
len(test_ds)
```

```
↳ 4
```

```
train_ds=train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds=val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
train_ds=test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
```

```
for image_batch,labels_batch in train_ds.take(1):
    print(image_batch.shape)
    print(image_batch[0].numpy()/255)
```

```

[32, 640, 640, 3]
[[[0.7372549 0.7647059 0.7882353 ]
 [0.7372549 0.7647059 0.7882353 ]
 [0.7372549 0.7647059 0.7882353 ]
 ...
 [0.3137255 0.31764707 0.19607843]
 [0.3019608 0.30588236 0.18039216]
 [0.2901961 0.29411766 0.16862746]]]

[[[0.7372549 0.7647059 0.7882353 ]
 [0.7372549 0.7647059 0.7882353 ]
 [0.7372549 0.7647059 0.7882353 ]
 ...
 [0.31764707 0.32156864 0.2
 [0.29411766 0.29803923 0.1764706 ]
 [0.27058825 0.27450982 0.14901961]]]

[[[0.7372549 0.7647059 0.7882353 ]
 [0.7372549 0.7647059 0.7882353 ]
 [0.7372549 0.7647059 0.7882353 ]
 ...
 [0.34117648 0.34117648 0.23137255]
 [0.3137255 0.3137255 0.20392157]
 [0.28235295 0.28627452 0.16470589]]]

...
[[[0.45882353 0.42745098 0.3764706 ]
 [0.4392157 0.40784314 0.35686275]
 [0.42745098 0.40392157 0.34901962]
 ...
 [0.8392157 0.5411765 0.40784314]
 [0.8352941 0.5411765 0.41568628]
 [0.8352941 0.5411765 0.41568628]]]

[[[0.45882353 0.42745098 0.3764706 ]
 [0.44313726 0.4117647 0.36078432]
 [0.43529412 0.4117647 0.35686275]
 ...
 [0.8392157 0.5411765 0.40784314]
 [0.8392157 0.54509807 0.41960785]
 [0.8392157 0.54509807 0.41960785]]]

[[[0.45490196 0.42352942 0.37254903]
 [0.44313726 0.4117647 0.36078432]
 [0.44313726 0.41960785 0.3647059]
 ...
 [0.84313726 0.54509807 0.4117647]
 [0.8392157 0.54509807 0.41960785]
 [0.8392157 0.54509807 0.41960785]]]

resize_and_rescale=tf.keras.Sequential([
    layers.Resizing(IMAGE_SIZE,IMAGE_SIZE),
    layers.Rescaling(1.0/255)
])

data_augmentation=tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2)
])

input_shape=(BATCH_SIZE,IMAGE_SIZE,IMAGE_SIZE,CHANNELS)
n_classes=3

model=models.Sequential([
    layers.Input(shape=(IMAGE_SIZE,IMAGE_SIZE,CHANNELS)),
    data_augmentation,
    resize_and_rescale,
    layers.Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=input_shape),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=input_shape),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),
    layers.Dense(64,activation='relu'),
])

```

```

layers.Dense(n_classes,activation='softmax'),
])
model.summary()

→ /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential_5"

```

Layer (type)	Output Shape	Param #
sequential_3 (Sequential)	(None, 640, 640, 3)	0
sequential_2 (Sequential)	(None, 640, 640, 3)	0
conv2d_1 (Conv2D)	(None, 638, 638, 32)	896
max_pooling2d (MaxPooling2D)	(None, 319, 319, 32)	0
conv2d_2 (Conv2D)	(None, 317, 317, 32)	9,248
max_pooling2d_1 (MaxPooling2D)	(None, 158, 158, 32)	0
conv2d_3 (Conv2D)	(None, 156, 156, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 78, 78, 64)	0
conv2d_4 (Conv2D)	(None, 76, 76, 64)	36,928
max_pooling2d_3 (MaxPooling2D)	(None, 38, 38, 64)	0
conv2d_5 (Conv2D)	(None, 36, 36, 64)	36,928
max_pooling2d_4 (MaxPooling2D)	(None, 18, 18, 64)	0
conv2d_6 (Conv2D)	(None, 16, 16, 64)	36,928
max_pooling2d_5 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_7 (Conv2D)	(None, 6, 6, 64)	36,928
max_pooling2d_6 (MaxPooling2D)	(None, 3, 3, 64)	0
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36,928
dense_1 (Dense)	(None, 3)	195

Total params: 213,475 (833.89 KB)

```

model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)
BATCH_SIZE

```

→ 32

```

history=model.fit(
    train_ds,
    batch_size=BATCH_SIZE,
    validation_data=val_ds,
    verbose=1,
    epochs=EPOCHS,
)

```

→

12/4/24, 1:05 AM

Animal disease identifier - Colab

```
4/4 ━━━━━━━━━━ 1s 374ms/step - accuracy: 0.6882 - loss: 0.7888 - val_accuracy: 0.6031 - val_loss: 0.8991
Epoch 32/50
4/4 ━━━━━━━━━━ 3s 372ms/step - accuracy: 0.6750 - loss: 0.8158 - val_accuracy: 0.7031 - val_loss: 0.8062
Epoch 33/50
4/4 ━━━━━━━━━━ 3s 373ms/step - accuracy: 0.6735 - loss: 0.8021 - val_accuracy: 0.7031 - val_loss: 0.8214
Epoch 34/50
4/4 ━━━━━━━━━━ 3s 373ms/step - accuracy: 0.6715 - loss: 0.8242 - val_accuracy: 0.7031 - val_loss: 0.8239
Epoch 35/50
4/4 ━━━━━━━━━━ 3s 380ms/step - accuracy: 0.6292 - loss: 0.8634 - val_accuracy: 0.6875 - val_loss: 0.8239
Epoch 36/50
4/4 ━━━━━━━━━━ 3s 375ms/step - accuracy: 0.6534 - loss: 0.8026 - val_accuracy: 0.6250 - val_loss: 0.8393
Epoch 37/50
4/4 ━━━━━━━━━━ 3s 376ms/step - accuracy: 0.6683 - loss: 0.7757 - val_accuracy: 0.6094 - val_loss: 0.8358
Epoch 38/50
4/4 ━━━━━━━━━━ 3s 371ms/step - accuracy: 0.6760 - loss: 0.7921 - val_accuracy: 0.6094 - val_loss: 0.8322
Epoch 39/50
4/4 ━━━━━━━━━━ 1s 372ms/step - accuracy: 0.6692 - loss: 0.7788 - val_accuracy: 0.6562 - val_loss: 0.8149
Epoch 40/50
4/4 ━━━━━━━━━━ 1s 372ms/step - accuracy: 0.6545 - loss: 0.8019 - val_accuracy: 0.7031 - val_loss: 0.8025
Epoch 41/50
4/4 ━━━━━━━━━━ 3s 378ms/step - accuracy: 0.6377 - loss: 0.8360 - val_accuracy: 0.6875 - val_loss: 0.8051
Epoch 42/50
4/4 ━━━━━━━━━━ 1s 371ms/step - accuracy: 0.6988 - loss: 0.7852 - val_accuracy: 0.6094 - val_loss: 0.8425
Epoch 43/50
4/4 ━━━━━━━━━━ 1s 370ms/step - accuracy: 0.6524 - loss: 0.7825 - val_accuracy: 0.6719 - val_loss: 0.7967
Epoch 44/50
4/4 ━━━━━━━━━━ 1s 371ms/step - accuracy: 0.6883 - loss: 0.7640 - val_accuracy: 0.6094 - val_loss: 0.8288
Epoch 45/50
4/4 ━━━━━━━━━━ 3s 372ms/step - accuracy: 0.7071 - loss: 0.7417 - val_accuracy: 0.6562 - val_loss: 0.8371
Epoch 46/50
4/4 ━━━━━━━━━━ 3s 370ms/step - accuracy: 0.7168 - loss: 0.7191 - val_accuracy: 0.6562 - val_loss: 0.8276
Epoch 47/50
4/4 ━━━━━━━━━━ 2s 409ms/step - accuracy: 0.6788 - loss: 0.7559 - val_accuracy: 0.6406 - val_loss: 0.8199
Epoch 48/50
4/4 ━━━━━━━━━━ 2s 374ms/step - accuracy: 0.6343 - loss: 0.8072 - val_accuracy: 0.6875 - val_loss: 0.8035
Epoch 49/50
4/4 ━━━━━━━━━━ 1s 374ms/step - accuracy: 0.6364 - loss: 0.8026 - val_accuracy: 0.6875 - val_loss: 0.7859
Epoch 50/50
4/4 ━━━━━━━━━━ 1s 373ms/step - accuracy: 0.7441 - loss: 0.7056 - val_accuracy: 0.5781 - val_loss: 0.8604
```

train_ds

```
↳ <_PrefetchDataset element_spec=(TensorSpec(shape=(None, 640, 640, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32, name=None))>
```

tf.__version__

```
↳ 2.17.1
```

scores=model.evaluate(test_ds)

```
↳ 4/4 ━━━━━━━━ 4s 108ms/step - accuracy: 0.5406 - loss: 1.0267
```

scores

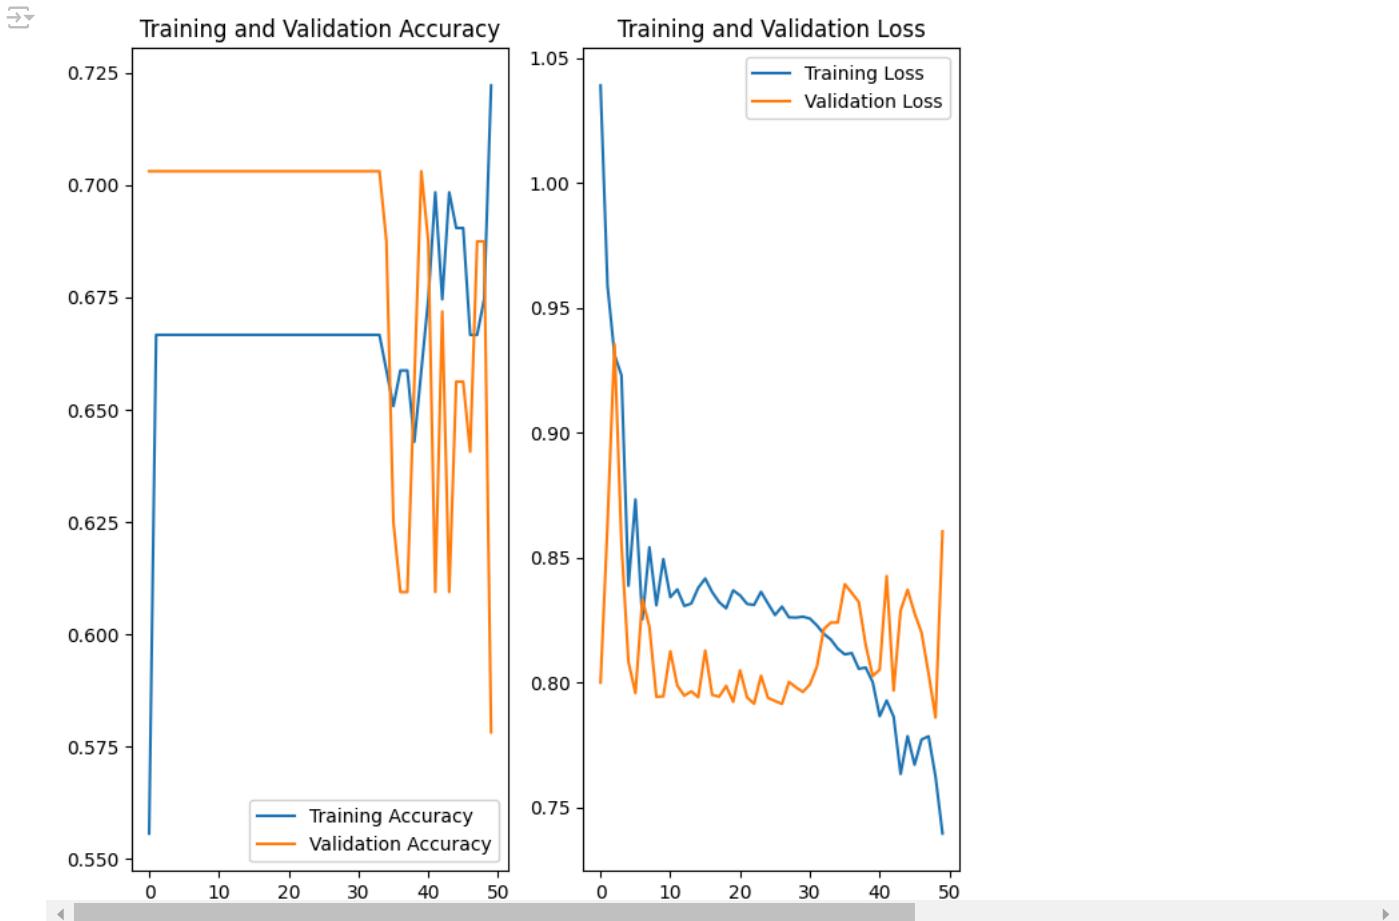
```
↳ [1.0516738891601562, 0.53125]
```

history.history.keys()

```
↳ dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

```
acc=history.history['accuracy']
val_acc=history.history['val_accuracy']
loss=history.history['loss']
val_loss=history.history['val_loss']
```

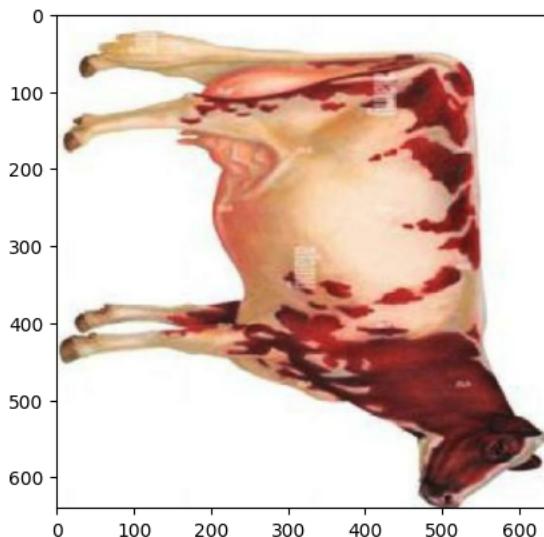
```
plt.figure(figsize=(8,8))
plt.subplot(1,2,1)
plt.plot(range(EPOCHS),acc,label='Training Accuracy')
plt.plot(range(EPOCHS),val_acc,label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.subplot(1,2,2)
plt.plot(range(EPOCHS),loss,label='Training Loss')
plt.plot(range(EPOCHS),val_loss,label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
import numpy as np

for images_batch,labels_batch in test_ds.take(1):
    first_image=images_batch[0].numpy().astype('uint8')
    first_label=labels_batch[0].numpy()
    print("first image to predict")
    plt.imshow(first_image)
    print("actual label:",class_names[first_label])
    batch_prediction=model.predict(images_batch)
    print("predicted label:",class_names[np.argmax(batch_prediction[0])])
```

→ first image to predict
actual label: train
1/1 0s 335ms/step
predicted label: train



```
def predict(model,img):
    img_array=tf.keras.preprocessing.image.img_to_array(images[i].numpy())
    img_array=tf.expand_dims(img_array,0)

    predictions=model.predict(img_array)
    predicted_class=class_names[np.argmax(predictions[0])]
```

```

confidence=round(100*(np.max(predictions[0])),2)
return predicted_class,confidence

plt.figure(figsize=(15,15))
for images,labels in test_ds.take(1):
    for i in range(9):
        ax=plt.subplot(3,3,i+1)
        plt.imshow(images[i].numpy().astype("uint8"))
        predicted_class,confidence=predict(model,images[i].numpy())
        actual_class=class_names[labels[i]]
        plt.title(f"Actual:{actual_class},\n Predicted:{predicted_class}.\n Confidence:{confidence}%")
        plt.axis("off")

```

1/1 1s 694ms/step
 1/1 0s 36ms/step
 1/1 0s 39ms/step
 1/1 0s 40ms/step
 1/1 0s 33ms/step
 1/1 0s 41ms/step
 1/1 0s 32ms/step
 1/1 0s 40ms/step
 1/1 0s 38ms/step

Actual:train,
 Predicted:train.
 Confidence:60.09%



Actual:valid,
 Predicted:train.
 Confidence:68.26%



Actual:valid,
 Predicted:valid.
 Confidence:82.83%



Actual:valid,
 Predicted:train.
 Confidence:75.17%



Actual:train,
 Predicted:train.
 Confidence:48.49%



Actual:train,
 Predicted:valid.
 Confidence:51.02%



Actual:train,
 Predicted:train.
 Confidence:64.52%



Actual:train,
 Predicted:train.
 Confidence:51.07%



Actual:train,
 Predicted:train.
 Confidence:64.37%



```

import os
save_dir="/content/drive/MyDrive/Disease dataset/models/"
model_version=1
os.makedirs(f"{save_dir}{model_version}",exist_ok=True)
model.save(f"{save_dir}{model_version}.keras")

model= models.Sequential([
    layers.Input(shape=(IMAGE_SIZE,IMAGE_SIZE,CHANNELS)),
    data_augmentation,
    resize_and_rescale,
    layers.Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=(128,128,3)),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=(128,128,3)),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,kernel_size=(3,3),activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),
    layers.Dense(64,activation='relu'),
    layers.Dense(n_classes,activation='softmax')
])

```

→ /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
model.summary()
```

→ Model: "sequential_7"

Layer (type)	Output Shape	Param #
sequential_3 (Sequential)	(None, 640, 640, 3)	0
sequential_2 (Sequential)	(None, 640, 640, 3)	0
conv2d_14 (Conv2D)	(None, 638, 638, 32)	896
max_pooling2d_13 (MaxPooling2D)	(None, 319, 319, 32)	0
conv2d_15 (Conv2D)	(None, 317, 317, 32)	9,248
max_pooling2d_14 (MaxPooling2D)	(None, 158, 158, 32)	0
conv2d_16 (Conv2D)	(None, 156, 156, 64)	18,496
max_pooling2d_15 (MaxPooling2D)	(None, 78, 78, 64)	0
conv2d_17 (Conv2D)	(None, 76, 76, 64)	36,928
max_pooling2d_16 (MaxPooling2D)	(None, 38, 38, 64)	0
conv2d_18 (Conv2D)	(None, 36, 36, 64)	36,928
max_pooling2d_17 (MaxPooling2D)	(None, 18, 18, 64)	0
conv2d_19 (Conv2D)	(None, 16, 16, 64)	36,928
max_pooling2d_18 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_2 (Flatten)	(None, 4096)	0
dense_4 (Dense)	(None, 64)	262,208
dense_5 (Dense)	(None, 3)	195

```

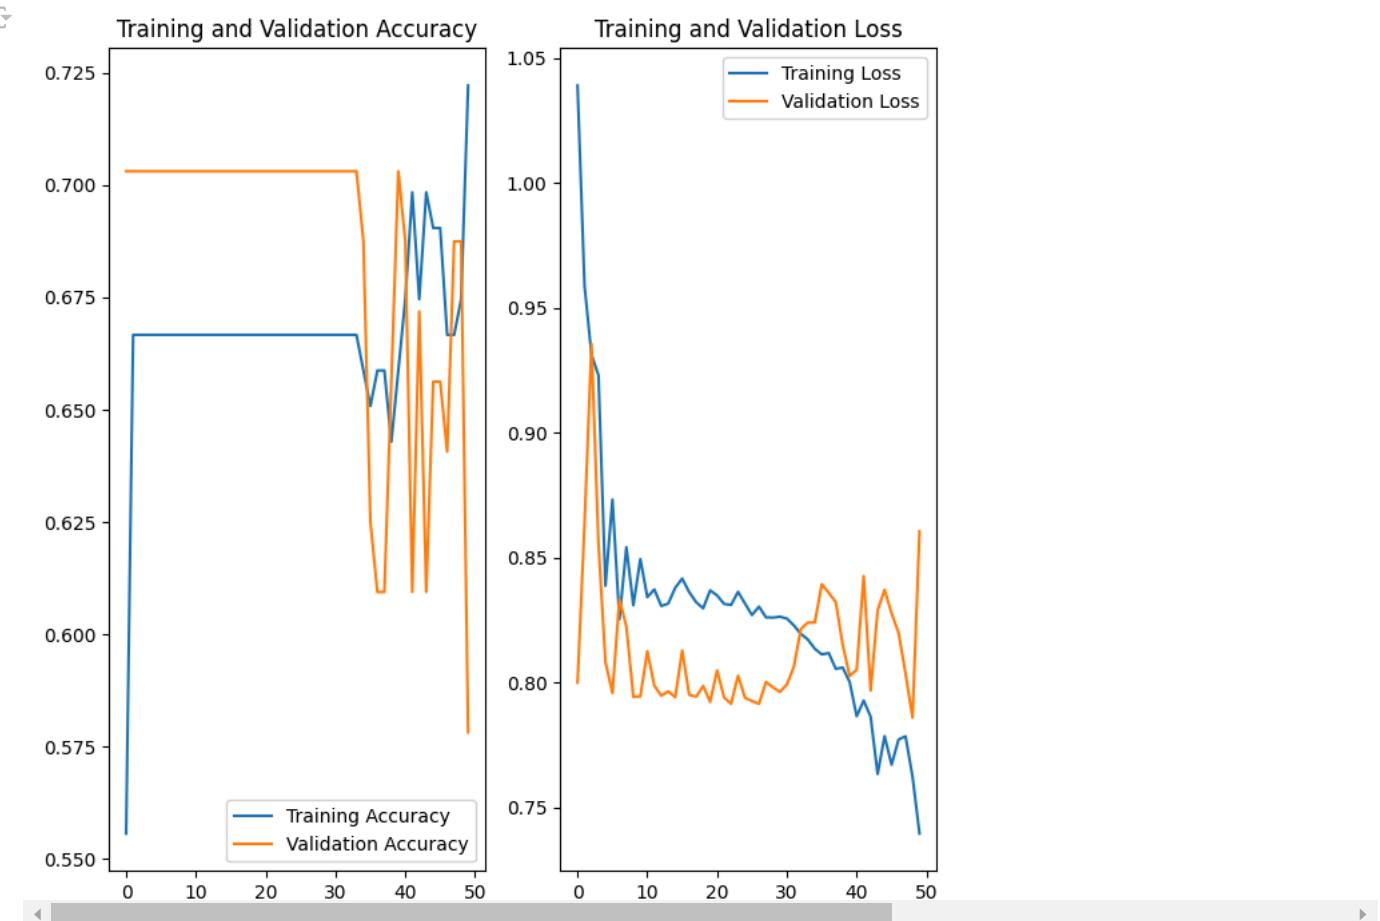
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)

history=model.fit(
    train_ds,
    batch_size=BATCH_SIZE,
    validation_data=val_ds,

```

```
verbose=1,  
epochs=EPOCHS,  
)  
→ Epoch 22/50  
4/4 ━━━━━━━━ 1s 377ms/step - accuracy: 0.6882 - loss: 0.7519 - val_accuracy: 0.6719 - val_loss: 0.8136  
Epoch 23/50  
4/4 ━━━━━━ 3s 376ms/step - accuracy: 0.6830 - loss: 0.7858 - val_accuracy: 0.6250 - val_loss: 0.8155  
Epoch 24/50  
4/4 ━━━━━━ 1s 374ms/step - accuracy: 0.6822 - loss: 0.7477 - val_accuracy: 0.5938 - val_loss: 0.8227  
Epoch 25/50  
4/4 ━━━━━━ 1s 375ms/step - accuracy: 0.6851 - loss: 0.7502 - val_accuracy: 0.5625 - val_loss: 0.8355  
Epoch 26/50  
4/4 ━━━━━━ 3s 413ms/step - accuracy: 0.6548 - loss: 0.8003 - val_accuracy: 0.5938 - val_loss: 0.8060  
Epoch 27/50  
4/4 ━━━━━━ 1s 378ms/step - accuracy: 0.6593 - loss: 0.7833 - val_accuracy: 0.6094 - val_loss: 0.7856  
Epoch 28/50  
4/4 ━━━━━━ 3s 410ms/step - accuracy: 0.7022 - loss: 0.7473 - val_accuracy: 0.5625 - val_loss: 0.8413  
Epoch 29/50  
4/4 ━━━━━━ 1s 378ms/step - accuracy: 0.6681 - loss: 0.7632 - val_accuracy: 0.5625 - val_loss: 0.8552  
Epoch 30/50  
4/4 ━━━━━━ 3s 373ms/step - accuracy: 0.6989 - loss: 0.7630 - val_accuracy: 0.6250 - val_loss: 0.7860  
Epoch 31/50  
4/4 ━━━━━━ 3s 372ms/step - accuracy: 0.6999 - loss: 0.7529 - val_accuracy: 0.5781 - val_loss: 0.8355  
Epoch 32/50  
4/4 ━━━━━━ 1s 373ms/step - accuracy: 0.6388 - loss: 0.8057 - val_accuracy: 0.6094 - val_loss: 0.8415  
Epoch 33/50  
4/4 ━━━━━━ 1s 373ms/step - accuracy: 0.7074 - loss: 0.7550 - val_accuracy: 0.6250 - val_loss: 0.8062  
Epoch 34/50  
4/4 ━━━━━━ 3s 370ms/step - accuracy: 0.7198 - loss: 0.6983 - val_accuracy: 0.5625 - val_loss: 0.8387  
Epoch 35/50  
4/4 ━━━━━━ 1s 375ms/step - accuracy: 0.6904 - loss: 0.7720 - val_accuracy: 0.6250 - val_loss: 0.7848  
Epoch 36/50  
4/4 ━━━━━━ 3s 388ms/step - accuracy: 0.6798 - loss: 0.7596 - val_accuracy: 0.6406 - val_loss: 0.8039  
Epoch 37/50  
4/4 ━━━━━━ 1s 372ms/step - accuracy: 0.7167 - loss: 0.7465 - val_accuracy: 0.6562 - val_loss: 0.8268  
Epoch 38/50  
4/4 ━━━━━━ 2s 406ms/step - accuracy: 0.6903 - loss: 0.7397 - val_accuracy: 0.6250 - val_loss: 0.8108  
Epoch 39/50  
4/4 ━━━━━━ 1s 372ms/step - accuracy: 0.6956 - loss: 0.7549 - val_accuracy: 0.6406 - val_loss: 0.7765  
Epoch 40/50  
4/4 ━━━━━━ 3s 371ms/step - accuracy: 0.6979 - loss: 0.7828 - val_accuracy: 0.6875 - val_loss: 0.7727  
Epoch 41/50  
4/4 ━━━━━━ 1s 371ms/step - accuracy: 0.6851 - loss: 0.7784 - val_accuracy: 0.6562 - val_loss: 0.7977  
Epoch 42/50  
4/4 ━━━━━━ 1s 373ms/step - accuracy: 0.7095 - loss: 0.7347 - val_accuracy: 0.6719 - val_loss: 0.8139  
Epoch 43/50  
4/4 ━━━━━━ 1s 377ms/step - accuracy: 0.7180 - loss: 0.7223 - val_accuracy: 0.6719 - val_loss: 0.8228  
Epoch 44/50  
4/4 ━━━━━━ 3s 373ms/step - accuracy: 0.6697 - loss: 0.7730 - val_accuracy: 0.6250 - val_loss: 0.8603  
Epoch 45/50  
4/4 ━━━━━━ 1s 373ms/step - accuracy: 0.7031 - loss: 0.7712 - val_accuracy: 0.6562 - val_loss: 0.8032  
Epoch 46/50  
4/4 ━━━━━━ 3s 373ms/step - accuracy: 0.7230 - loss: 0.7456 - val_accuracy: 0.6719 - val_loss: 0.8039  
Epoch 47/50  
4/4 ━━━━━━ 1s 374ms/step - accuracy: 0.6915 - loss: 0.7531 - val_accuracy: 0.6562 - val_loss: 0.7870  
Epoch 48/50  
4/4 ━━━━━━ 3s 371ms/step - accuracy: 0.7000 - loss: 0.6877 - val_accuracy: 0.6406 - val_loss: 0.8369  
Epoch 49/50  
4/4 ━━━━━━ 1s 371ms/step - accuracy: 0.7295 - loss: 0.7351 - val_accuracy: 0.6562 - val_loss: 0.7860  
Epoch 50/50  
4/4 ━━━━━━ 2s 436ms/step - accuracy: 0.7269 - loss: 0.7457 - val_accuracy: 0.6719 - val_loss: 0.7849
```

```
plt.figure(figsize=(8,8))  
plt.subplot(1,2,1)  
plt.plot(range(EPOCHS),acc,label='Training Accuracy')  
plt.plot(range(EPOCHS),val_acc,label='Validation Accuracy')  
plt.legend(loc='lower right')  
plt.title('Training and Validation Accuracy')  
  
plt.subplot(1,2,2)  
plt.plot(range(EPOCHS),loss,label='Training Loss')  
plt.plot(range(EPOCHS),val_loss,label='Validation Loss')  
plt.legend(loc='upper right')  
plt.title('Training and Validation Loss')  
plt.show()
```



```
plt.figure(figsize=(15,15))
for images,labels in test_ds.take(1):
    for i in range(9):
        ax=plt.subplot(3,3,i+1)
        plt.imshow(images[i].numpy().astype("uint8"))
        predicted_class,confidence=predict(model,images[i].numpy())
        actual_class=class_names[labels[i]]
        plt.title(f"Actual:{actual_class},\n Predicted:{predicted_class}.\n Confidence:{confidence}%")
        plt.axis("off")
```

```
1/1 ━━━━━━ 0s 217ms/step
1/1 ━━━━ 0s 37ms/step
1/1 ━━ 0s 28ms/step
1/1 ━ 0s 27ms/step
1/1 0s 31ms/step
1/1 0s 28ms/step
1/1 0s 35ms/step
1/1 0s 28ms/step
1/1 0s 30ms/step
```

Actual:train,
Predicted:train.
Confidence:56.95%



Actual:test,
Predicted:train.
Confidence:69.46%



Actual:valid,
Predicted:train.
Confidence:63.89%



Actual:test,
Predicted:train.
Confidence:65.87%



Actual:train,
Predicted:train.
Confidence:47.4%



Actual:train,
Predicted:train.
Confidence:81.83%



Actual:train,
Predicted:train.
Confidence:69.36%



Actual:valid,
Predicted:train.
Confidence:79.19%



Actual:valid,
Predicted:train.
Confidence:68.52%



```
import os
save_dir="/content/drive/MyDrive/Disease dataset/models/"
model_version=2
os.makedirs(f"{save_dir}{model_version}",exist_ok=True)
model.save(f"{save_dir}{model_version}.keras")
```

```
model=models.Sequential([
    layers.Input(shape=(IMAGE_SIZE,IMAGE_SIZE,CHANNELS)),
    data_augmentation,
    resize_and_rescale,
    layers.Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=(640,640,3))
```

```

        layers.MaxPooling2D((2,2)),
        layers.Flatten(),
        layers.Dense(64,activation='relu'),
        layers.Dense(n_classes,activation='softmax')
    ])
model.summary()

```

→ /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_.onv.py:107: UserWarning: Do not pass an `input_shape`/`super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential_8"

Layer (type)	Output Shape	Param #
sequential_3 (Sequential)	(None, 640, 640, 3)	0
sequential_2 (Sequential)	(None, 640, 640, 3)	0
conv2d_20 (Conv2D)	(None, 638, 638, 32)	896
max_pooling2d_19 (MaxPooling2D)	(None, 319, 319, 32)	0
flatten_3 (Flatten)	(None, 3256352)	0
dense_6 (Dense)	(None, 64)	208,406,592
dense_7 (Dense)	(None, 3)	195

Total params: 208,407,683 (795.01 MB)

Trainable params: 208,407,683 (795.01 MB)

```

model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)

```

```

history=model.fit(
    train_ds,
    batch_size=BATCH_SIZE,
    validation_data=val_ds,
    verbose=1,
    epochs=EPOCHS,
)

```

→ Epoch 22/50
 4/4 1s 341ms/step - accuracy: 0.6967 - loss: 0.7488 - val_accuracy: 0.6094 - val_loss: 0.8848
 Epoch 23/50
 4/4 1s 341ms/step - accuracy: 0.6978 - loss: 0.7338 - val_accuracy: 0.6719 - val_loss: 0.7929
 Epoch 24/50
 4/4 3s 339ms/step - accuracy: 0.6884 - loss: 0.7037 - val_accuracy: 0.6406 - val_loss: 0.8243
 Epoch 25/50
 4/4 3s 341ms/step - accuracy: 0.7043 - loss: 0.7046 - val_accuracy: 0.6406 - val_loss: 0.8670
 Epoch 26/50
 4/4 3s 337ms/step - accuracy: 0.6875 - loss: 0.6991 - val_accuracy: 0.6094 - val_loss: 0.9327
 Epoch 27/50
 4/4 3s 337ms/step - accuracy: 0.6677 - loss: 0.7140 - val_accuracy: 0.6250 - val_loss: 0.8643
 Epoch 28/50
 4/4 3s 342ms/step - accuracy: 0.7020 - loss: 0.6891 - val_accuracy: 0.6406 - val_loss: 0.8224
 Epoch 29/50
 4/4 1s 344ms/step - accuracy: 0.6710 - loss: 0.7760 - val_accuracy: 0.6250 - val_loss: 0.8907
 Epoch 30/50
 4/4 1s 341ms/step - accuracy: 0.7080 - loss: 0.7264 - val_accuracy: 0.6406 - val_loss: 0.8757
 Epoch 31/50
 4/4 3s 393ms/step - accuracy: 0.7127 - loss: 0.7101 - val_accuracy: 0.6094 - val_loss: 0.9244
 Epoch 32/50
 4/4 1s 339ms/step - accuracy: 0.7013 - loss: 0.7313 - val_accuracy: 0.6250 - val_loss: 0.9024
 Epoch 33/50
 4/4 3s 339ms/step - accuracy: 0.7033 - loss: 0.7308 - val_accuracy: 0.6406 - val_loss: 0.8483
 Epoch 34/50
 4/4 1s 344ms/step - accuracy: 0.6862 - loss: 0.7602 - val_accuracy: 0.6094 - val_loss: 0.9511
 Epoch 35/50